



CISC3024 Project

Classification of Street View House
Numbers Using PyTorch

AgustinYan DC125714



SVHN

Street View House Number dataset

- Real-world image dataset
- Similar to MNIST
- But it is harder and unsolved

Data Processing and Augmentation

```
class SVHNAugmentedDataset(torch.utils.data.Dataset):
    def __init__(self, root, split='train', download=True,
                 max_rotation=15, min_crop_size=16, max_aspect_ratio_change=0.1, train=True):
        self.dataset = SVHN(root=root, split=split, download=download) # load the dataset
        self.train = train # to control the train and test

        self.augment = A.Compose([ # when in train mode, do the augmentation
            A.Rotate(limit=max_rotation, p=0.5),
            A.RandomResizedCrop(height=32, width=32, scale=(min_crop_size/32, 1.0), ratio=(1-max_aspect_ratio_change, 1+max_aspect_ratio_change)),
            A.Normalize(mean=[0.4377, 0.4438, 0.4728], std=[0.1980, 0.2010, 0.1970]),
            ToTensorV2()
        ])

        self.normalize = A.Compose([ # when in test mode, only do normalization
            A.Normalize(mean=[0.4377, 0.4438, 0.4728], std=[0.1980, 0.2010, 0.1970]),
            ToTensorV2()
        ])

    def __len__(self):
        return len(self.dataset)

    def __getitem__(self, idx):
        image, label = self.dataset[idx]
        image = np.array(image)

        if self.train:
            augmented = self.augment(image=image)
        else:
            augmented = self.normalize(image=image)

        image = augmented['image']
        return image, label
```

Data Processing and Augmentation

```
# load and transform the dataset
train_dataset = SVHNAugmentedDataset(root='./data', split='train', download=True, train=True)
test_dataset = SVHNAugmentedDataset(root='./data', split='test', download=True, train=False)

# create dataloader
train_loader = DataLoader(train_dataset, batch_size=64, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=64, shuffle=False)
```

Neural Network Setup

```
class SmallVGG(nn.Module):
    def __init__(self):
        super(SmallVGG, self).__init__()
        self.conv_layers = nn.Sequential(
            nn.Conv2d(3, 8, kernel_size=3, padding=1), nn.GroupNorm(2, 8), nn.ReLU(),
            nn.Conv2d(8, 16, kernel_size=3, padding=1), nn.GroupNorm(4, 16), nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2), nn.Dropout(0.25),

            nn.Conv2d(16, 32, kernel_size=3, padding=1), nn.GroupNorm(8, 32), nn.ReLU(),
            nn.Conv2d(32, 32, kernel_size=3, padding=1), nn.GroupNorm(8, 32), nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2), nn.Dropout(0.25),

            nn.Conv2d(32, 32, kernel_size=3, padding=1), nn.GroupNorm(8, 32), nn.ReLU(),
            nn.Conv2d(32, 32, kernel_size=3, padding=1), nn.GroupNorm(8, 32), nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2), nn.Dropout(0.25)
        )

        self.fc_layers = nn.Sequential(
            nn.Linear(32 * 4 * 4, 256), nn.ReLU(),
            nn.Dropout(0.5),
            nn.Linear(256, 10) #output layer 10 classes
        )

    def forward(self, x):
        x = self.conv_layers(x)
        x = x.view(x.size(0), -1)
        x = self.fc_layers(x)
        return x
```

Training and Evaluation process

```
def train_and_evaluate(model, train_loader, test_loader, criterion, optimizer, num_epochs):
    train_losses = []
    test_losses = []
    test_accuracies = []
    all_labels = []
    all_scores = []

    for epoch in range(num_epochs):
        model.train()
        running_loss = 0.0

        # Training loop
        for images, labels in tqdm(train_loader):
            images, labels = images.to(device), labels.to(device)

            optimizer.zero_grad()
            outputs = model(images)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

            running_loss += loss.item() * images.size(0)

        epoch_train_loss = running_loss / len(train_loader.dataset)
        train_losses.append(epoch_train_loss)

        model.eval()
        test_loss = 0.0
        correct = 0
        total = 0
```

Training and Evaluation process

```
# Evaluation loop
with torch.no_grad():
    for images, labels in test_loader:
        images, labels = images.to(device), labels.to(device)
        outputs = model(images)
        loss = criterion(outputs, labels)
        test_loss += loss.item() * images.size(0)

    _, predicted = torch.max(outputs, 1)
    total += labels.size(0)
    correct += (predicted == labels).sum().item()

    all_labels.append(labels.cpu().numpy())
    all_scores.append(outputs.cpu().numpy())


epoch_test_loss = test_loss / len(test_loader.dataset)
test_losses.append(epoch_test_loss)
accuracy = 100 * correct / total
test_accuracies.append(accuracy)

print(f"Epoch [{epoch + 1}/{num_epochs}] - "
      f"Train Loss: {epoch_train_loss:.4f}, "
      f"Test Loss: {epoch_test_loss:.4f}, "
      f"Accuracy: {accuracy:.2f}%")


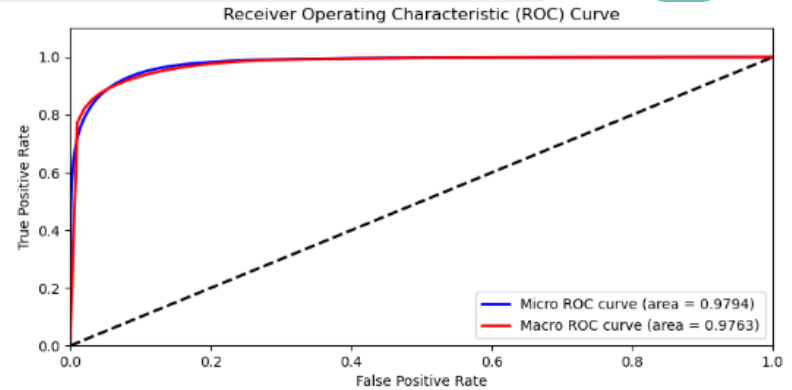

# Convert lists to numpy arrays for ROC AUC computation
all_labels = np.concatenate(all_labels)
all_scores = np.concatenate(all_scores)

return train_losses, test_losses, test_accuracies, all_labels, all_scores
```

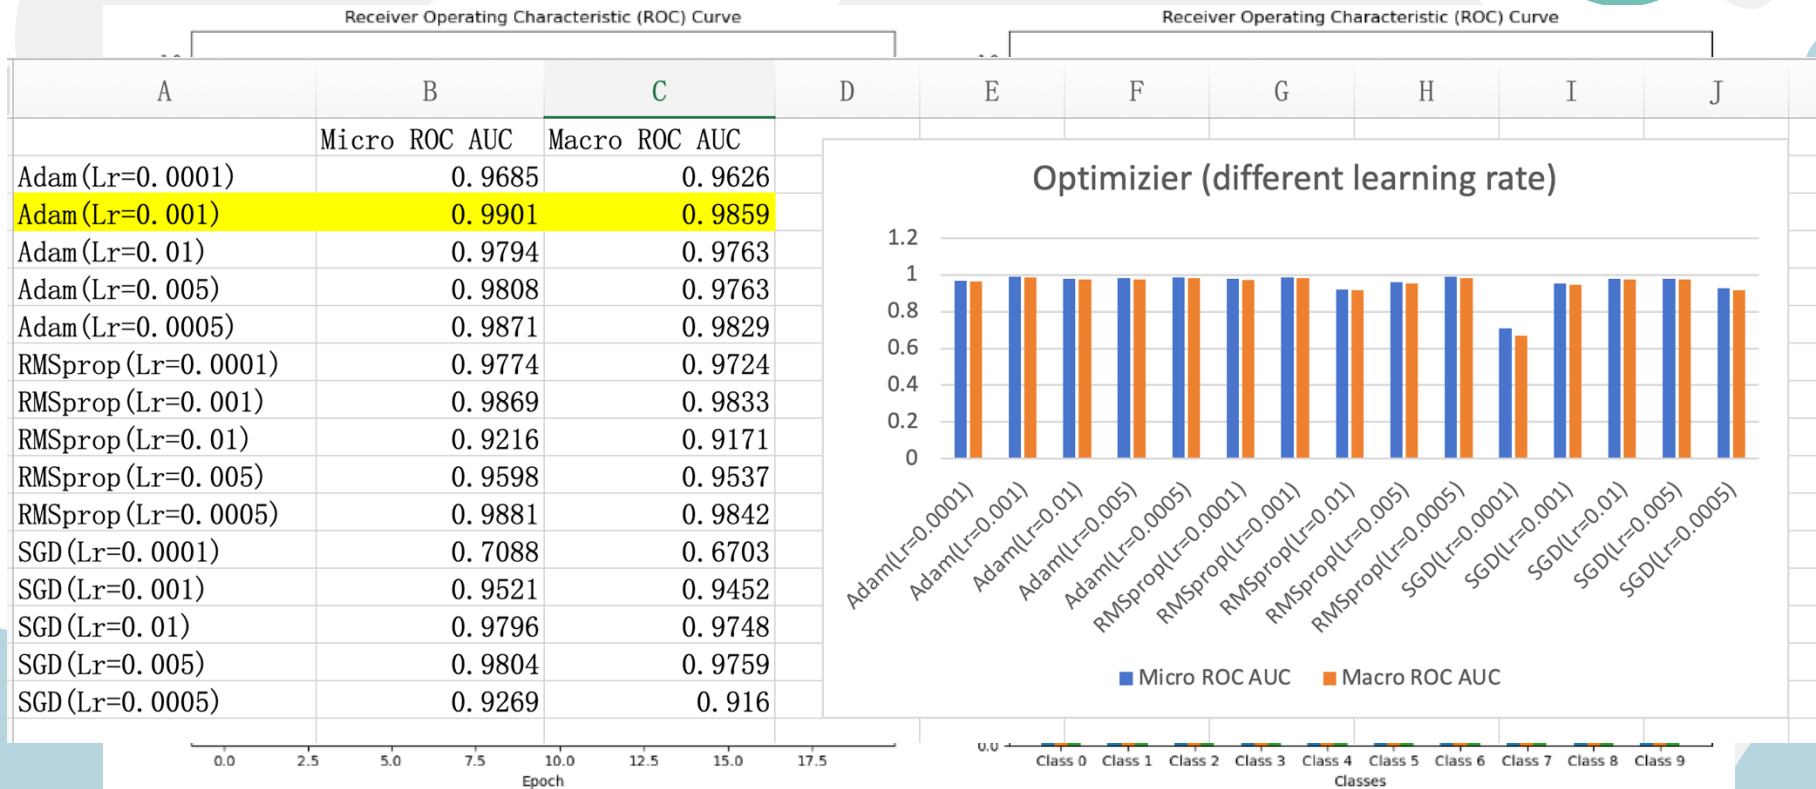
Analysis Training and Evaluation process



Optimizer (with different Learning rate)

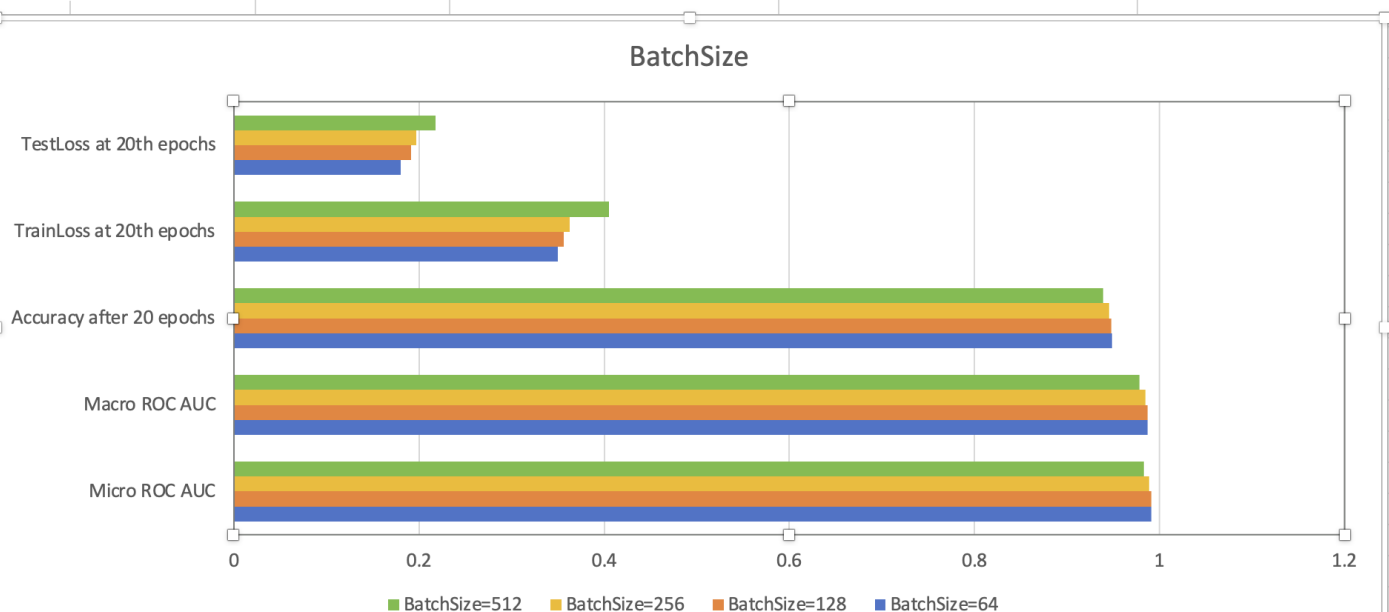


Optimizer (with different Learning rate)



BatchSize

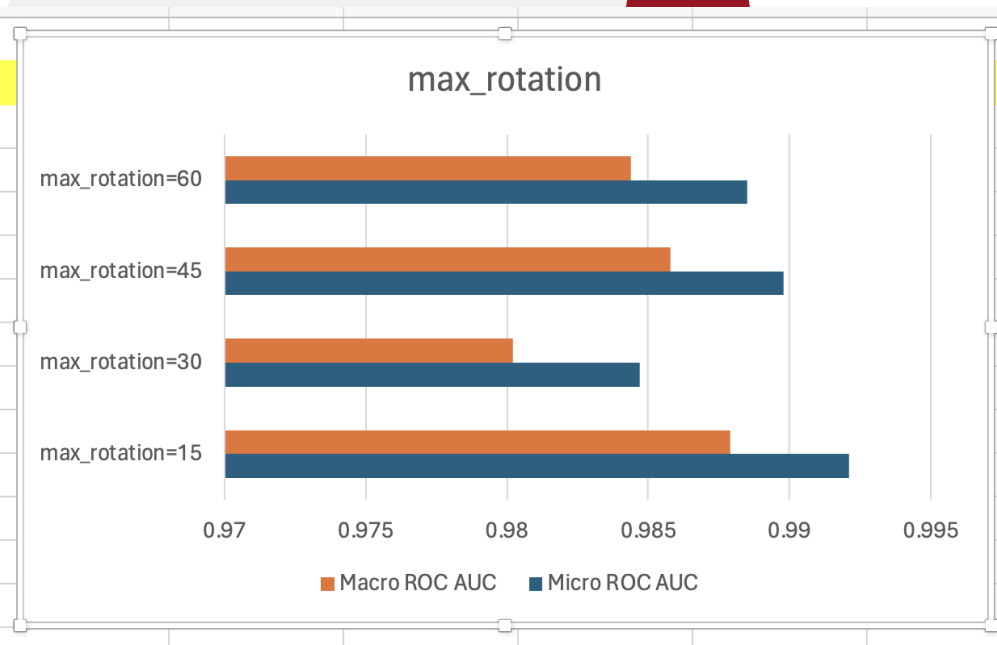
	Micro ROC AUC	Macro ROC AUC	Accuracy after 20 epochs	TrainLoss at 20th epochs	TestLoss at 20th epochs
BatchSize=64	0.9913	0.9873	94.91%	0.35	0.18
BatchSize=128	0.991	0.987	94.79%	0.3564	0.1912
BatchSize=256	0.9889	0.9847	94.54%	0.363	0.1968
BatchSize=512	0.9832	0.9786	93.91%	0.4047	0.2175







Maximum angle of rotation

	Micro ROC AUC	Macro ROC AUC
max_rotation=15	0.9921	0.9879
max_rotation=30	0.9847	0.9802
max_rotation=45	0.9898	0.9858
max_rotation=60	0.9885	0.9844

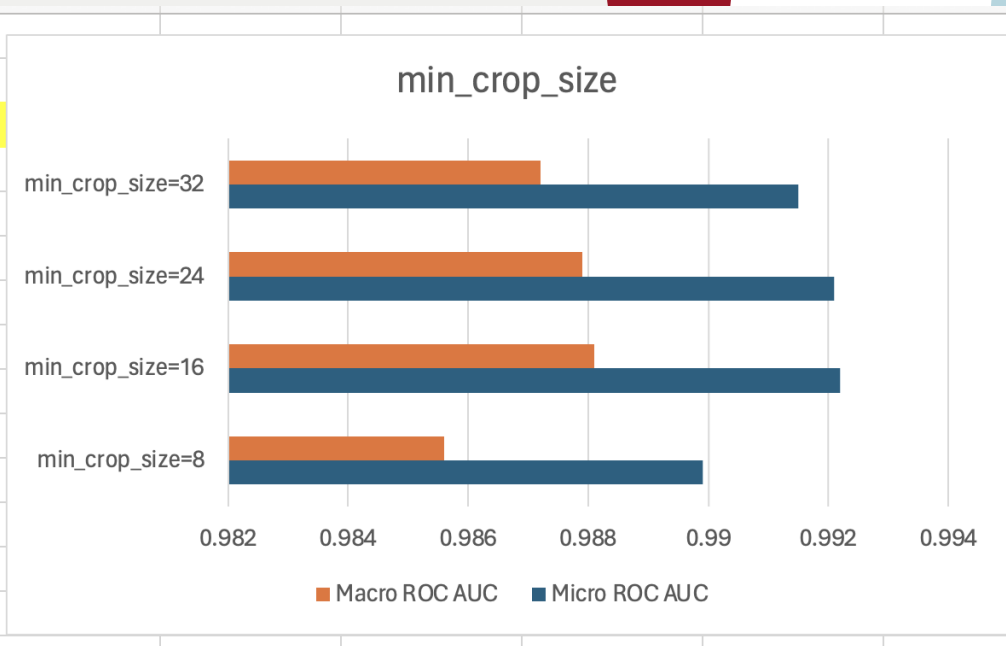
图表区







Minimum Crop Size



Minimum Crop Size

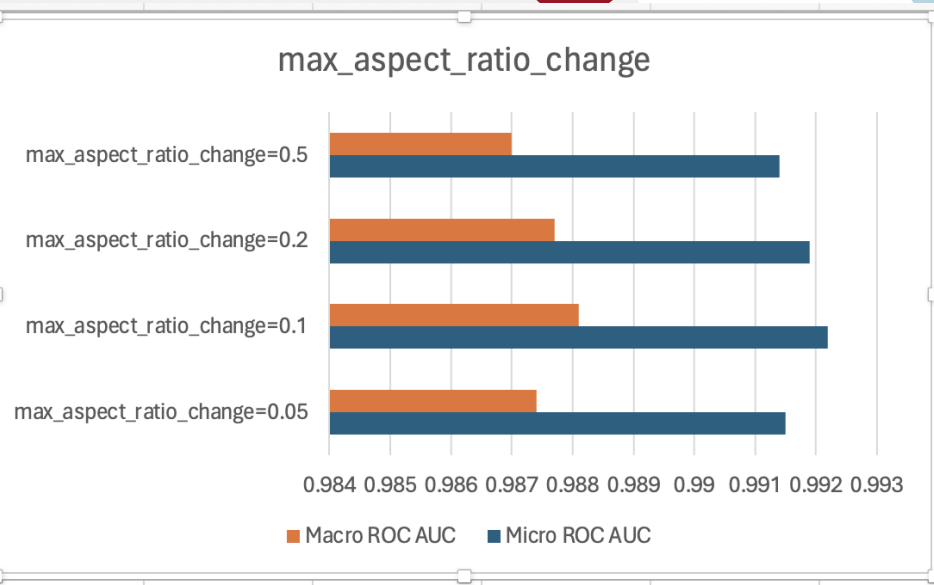
[illegible]

Maximum aspect ratio change

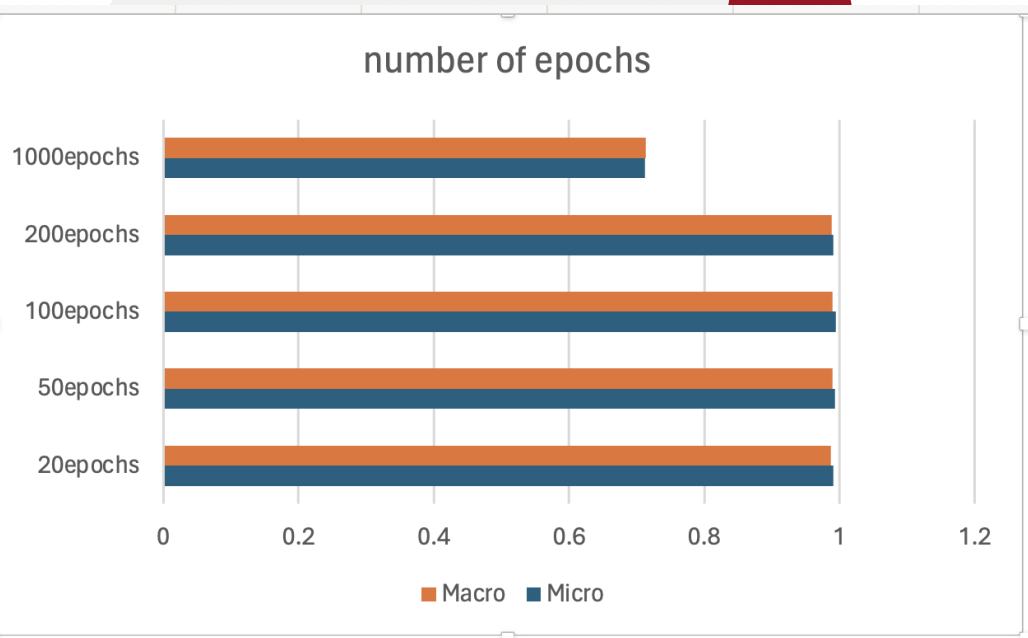









Maximum aspect ratio change

	Micro ROC AUC	Macro ROC AUC
max_aspect_ratio_change=0.05	0.9915	0.9874
max_aspect_ratio_change=0.1	0.9922	0.9881
max_aspect_ratio_change=0.2	0.9919	0.9877
max_aspect_ratio_change=0.5	0.9914	0.987





Number of Epochs

[illegible]



100 Epochs and
1000 Epochs

Final performance of the model



The background is a light blue gradient with various organic, blob-like shapes in orange, red, yellow, and teal. A large white circle is centered on the left side, containing the text.

Model Demonstration

With Sketchpad in Gradio



Thanks!

For your kind listening

Any questions?